

KIMI K3 VS CLAUDE FABLE 5

The strongest open-weight model ever released is second-best at nearly everything — and first at exactly one thing.

Published third-party figures only · empty rows where no comparable number exists

VERDICT

- **Claude Fable 5 leads overall.** Top of the aggregate indices, first on every published software-engineering benchmark, first on long-horizon agentic work.
- **Kimi K3 wins frontend.** Ranked #1 on Arena.ai's Frontend Code arena — above Fable 5. The only major leaderboard where it finishes ahead.
- **K3 wins on economics.** 30% of Fable 5's list price on both sides of the meter, roughly half Opus 4.8's cost per task, and open weights.
- **No head-to-head coding comparison exists.** K3 has not been scored on the harnesses where Fable 5 posts its numbers. Anyone presenting one is estimating.

Moonshot AI shipped Kimi K3 on 16 July 2026: a 2.8-trillion-parameter mixture-of-experts model with a one-million-token context window, native vision, and full weights dated 27 July. It is the largest open-weight model ever released and the first Chinese model to compete with the top US systems on capability rather than on price alone. It beat Claude Opus 4.8 on general capability at roughly half the cost per task. It did not beat Claude Fable 5.

01 WHAT K3 ACTUALLY IS


Moonshot describes K3 as a mixture-of-experts model activating 16 of 896 experts, using Kimi Delta Attention, Attention Residuals, and a Stable LatentMoE design. The launch materials specify MXFP4 weights with MXFP8 activations and recommend a supernode of at least 64 accelerators for deployment — a reminder that "open weights" and "runnable by you" are different claims at this scale.

The comparison against Fable 5, Anthropic's generally available Mythos-class model released 9 June 2026:

Released	9 Jun 2026	16 Jul 2026
Architecture	not disclosed	MoE · 2.8T total · 16 of 896 active
Context window	1,000,000	1,048,576
Max output per request	128,000	not stated
Modalities in	text · image · file	text · image
Price / 1M input	\$10.00	\$3.00
Price / 1M output	\$50.00	\$15.00
Batch pricing	\$5 / \$25	—
Weights	proprietary	open, dated 27 Jul 2026

Two things stand out before a single benchmark. The context windows are effectively identical, so context is not a differentiator this generation — the meaningful number is what it costs to fill one, which is \$10 against \$3. And the open-weights story carries a caveat: at launch the weights were not public, which is why Artificial Analysis classified K3 as proprietary while Moonshot marketed it as the first open 3T-class model. Everything downstream of that depends on the 27 July release landing as promised.

02 GENERAL CAPABILITY



The ordering repeats across every aggregate index. K3 sits third on Artificial Analysis's Intelligence Index, behind only Fable 5 and GPT-5.6 Sol Max — while costing less than either. It places second on the Vals AI index. Nathan Lambert's framing is the fair one: Moonshot is going toe to toe with Anthropic and OpenAI on a fraction of the resources, and this is unambiguously the strongest open model yet released.


But "strongest open model" and "strongest model" are different claims, and only the first survives the data.

03 SOFTWARE ENGINEERING: THE COMPARISON NOBODY CAN MAKE

This is where most comparison pieces quietly invent numbers. There are none to invent from.

CODING BENCHMARKS — % RESOLVED

Moonshot has not published K3 results on these harnesses. Artificial Analysis has not yet assigned it weighted coding coverage.



The empty rows are the finding, not an omission.

Moonshot's own launch materials report self-run evaluations placing K3 ahead of Opus 4.8 max and GPT-5.5 high while losing to Fable 5 and GPT-5.6 Sol — but not on these harnesses, and self-reported. So the honest statement is narrow: **Fable 5 leads on every published software-engineering benchmark, and K3 has not been measured on them.** That is not the same as winning a contest, and anyone presenting it as one is overreaching. It is also the only thing the current evidence supports.

One caveat on the Fable 5 figures themselves. SWE-bench Verified at 95% is near saturation — Vals runs it on a minimal bash-tool-only agent harness so the burden falls on the model rather than the scaffold, but at that level the gap between a great model and the benchmark ceiling compresses to nothing and the score stops discriminating. FrontierCode Diamond at 29.3% is the more interesting number precisely because it is low: harder, less saturated evaluations carry more signal about production behaviour than a saturated headline does.

04 WHERE K3 WINS OUTRIGHT

ARENA.AI — FRONTEND CODE ARENA

Human preference ranking on interface generation. Public Elo not disclosed; shown as standing.

1 Kimi K3

ABOVE FABLE 5


2 Claude Fable 5

The only major leaderboard where K3 finishes ahead measures the thing you look at and judge: layout, components, the visual result.



Be precise about what this measures. Arena rankings are human-preference Elo on generated interface code — not algorithmic correctness, not repository navigation, not multi-file refactors. If your workload is screens, landing pages, component libraries or design-adjacent frontend work, the current leader is the cheaper model. That is an unusual sentence to be able to write about an open-weight release, and it is the single most commercially relevant finding in this comparison.

Two qualifications keep it honest. Human-preference Elo rewards output that *presents* well — a real signal for interface work, a weaker one than a functional test. And Fable 5 holds #1 on FrontierCode, a considerably harder and more correctness-oriented evaluation. Both facts are true simultaneously: K3 produces the better-looking interface, Fable 5 solves the harder engineering problem.

05 LONG-HORIZON AGENTIC WORK



Second place behind Fable 5 — and the jump from K2.6 is the most important number in this entire comparison, because it isn't about standing at all. It is about slope. A lab that moves 732 Elo points between releases is not a lab you assess on where it currently sits. Moonshot also reports K3 using 21% fewer output tokens than K2.6 on the Artificial Analysis index, which means the efficiency gain compounds with the capability gain rather than trading against it.



K3 does second-best work at roughly half the cost per task of Opus 4.8 — a model it finished ahead of on general capability. Fable 5's batch pricing halves its list rate to \$5/\$25, which narrows the gap for asynchronous workloads but does not close it.

06 THE OPEN-WEIGHTS QUESTION

K3 takes the open-weight crown from DeepSeek's 1.6T v4 Pro, close to doubling the parameter count of the largest previously available open model. Moonshot rounds 2.8 up and calls it the first open 3T-class release, which is marketing but not by much.


What that means in practice is narrower than the headline implies. MXFP4 weights, MXFP8 activations and a recommended supernode of at least 64 accelerators put self-hosting outside the reach of everyone except well-capitalised infrastructure teams. Nobody is running this on a workstation. The constituency for whom the weights genuinely change something is specific: sovereign and regulated deployments, organisations that cannot route data through a US API, research groups that need to inspect a frontier-scale model rather than query one, and teams fine-tuning for a domain the hosted models will not serve.

The second consequence is competitive rather than technical. An open model sitting second or third on the aggregate indices compresses the window in which any proprietary release stays uncontested. That window has been shortening release over release, and this is the first time it has closed to within one generation from the open side.

There is also an unresolved classification tension worth naming. On 17 July, Artificial Analysis listed K3 as proprietary because the weights were not public; Moonshot described it as open. Both were accurate on that date. Until the weights and the promised technical report are published, developers can evaluate the hosted model and inspect the open Kimi Code harness, but cannot reproduce weight-level serving or verify the training claims. That is why 27 July is the date to watch rather than 16 July — the release that matters has not happened yet.

07 READ EVERY NUMBER ABOVE WITH THIS IN MIND

Benchmark scores are not model properties. They are model-plus-harness properties, and the harness moves them further than most cross-vendor gaps do.



Two implications follow. First, any cross-model gap under roughly five points on a coding benchmark is uninterpretable unless the harness was held constant — which is most of the comparisons currently circulating. Second, if you are selecting a model for agentic work, the scaffold you wrap it in is a lever of comparable size to the model choice, and a considerably cheaper one to pull. Search and context retrieval are usually where the difference comes from.

This is also why K3's missing coding rows matter more than they first appear. K3 ships with an open Kimi Code harness — a genuine asset for reproducibility — but it means its available coding evidence is entangled with its own scaffold in a way that Fable 5's minimal-agent SWE-bench figures are not. Until K3 is scored on a standardised harness, the comparison stays unavailable rather than merely unfavourable.

08 THREE CAVEATS THAT BELONG IN ANY RECOMMENDATION

Fable 5 routes around its own benchmarks in guarded domains

Requests touching cybersecurity, biology, chemistry or model distillation fall back to Claude Opus 4.8 — Anthropic reports this fires in under 5% of sessions, and Fable pricing is not charged when it does. The practical consequence is that in those domains you receive Opus-tier capability regardless of which model you called. If your work lives in security research or computational biology, the Fable 5 benchmark table is not the table that applies to you.

K3's hallucination figure needs watching

One secondary source reports K3's raw accuracy rising to roughly 46% from 33% for K2.6, while its measured hallucination rate rose to around 51%. That figure comes from a single lower-quality source and is not corroborated in the primary evaluations, so it should be treated as unverified — but it points at a pattern worth testing on your own workload before trusting K3 for factual output without a verification step.

Anthropic's own documents disagree by four points

Fable 5's Terminal-Bench 2.1 result appears as 84.3% in the system card and 88.0% in a separate published table. The lower figure is used throughout this article. When a vendor's own artifacts disagree by four points on the same benchmark, that is a standing argument for running your own evaluation rather than procuring on a leaderboard.

09 WHICH TO USE

Claude Fable 5 when the task is long-horizon and autonomous, when correctness outweighs cost, when the work is hard multi-file engineering, or when you need file inputs and a 128K output ceiling in a single request. It leads every aggregate index and every published coding benchmark. It is also 3.3× the price on both sides of the meter, which is the whole of the argument against it.

Kimi K3 when the output is an interface — where it is currently the outright leader — or when volume makes price the binding constraint. Second-best across the board at 30% of list price and roughly half Opus 4.8's cost per task is a strong position. The open weights also make it the only one of the two available for on-premise deployment, fine-tuning, or any environment the model cannot leave. Verify factual output.

A practical routing rule if you have access to both: default to K3 for interface generation and high-volume drafting, escalate to Fable 5 for long autonomous runs and hard multi-file engineering, and re-examine whether Opus 4.8 still earns a slot in between. On this evidence it is beaten by K3 on general capability while costing nearly twice as much per task — the quiet finding in this comparison is not that an open model closed on the frontier, but that it passed a proprietary model still being sold at a premium.

THE TRAJECTORY IS THE ACTUAL STORY

Fable 5 leads today, and the lead is smaller than it was a generation ago. K3 moved 732 Elo points on long-horizon work in one release, from a lab operating on a fraction of the compute and capital available to Anthropic or OpenAI. Two markers decide how this reads by autumn: whether the weights publish on 27 July as dated, and whether K3 earns weighted coding coverage on standardised harnesses rather than harness-specific rows.

Until then the summary is the boring one, which is usually the true one. **Fable 5 is the better model. K3 is the better deal. And the gap is closing from below.**

Sources – Artificial Analysis (Intelligence Index; private long-horizon evaluation; cost per task) · Arena.ai Frontend Code arena · Vals AI · Anthropic Claude Fable 5 / Mythos 5 system card and published benchmark table · Moonshot AI launch materials · OpenRouter model pages and pricing · Scale SEAL public and private leaderboards · Morph internal SWE-bench Pro benchmarks · release coverage from Simon Willison, Nathan Lambert (Interconnects) and VentureBeat.

Method – Only published third-party figures are charted. Where no comparable measurement exists, the row is left visibly empty rather than estimated or inferred. GDPval-AA v2 results are Max configurations. Elo axes are truncated where noted to show separation; percentage axes run 0–100. Figures current as of 22 July 2026 and subject to change as evaluation coverage expands.